

EECE 494

Real-Time Systems Design

Sathish Gopalakrishnan
University of British Columbia



Course information

- Course website: <http://courses.ece.ubc.ca/494/>
- Announcements will be posted on the website - you should check the website frequently
- For some announcements, a course mailing list will be used
 - Instructions to sign up for the mailing list are on the website
- Instructor's office: 4045 Kaiser Building (Phone: 604-827-4343)
- Instructor's office hours
 - Mondays (2 p.m. - 3 p.m.) & Thursdays (11 a.m. to 12 noon)
 - Office hours start from the week of January 11
- Teaching Assistant: Bader Na'eem Al-Ahmad


Reading material

- No required textbook
- All required reading material will be available on the course website
 - You will need to use the UBC VPN to access (almost all) the readings
- Recommended material
 - [Hard Real-Time Computing Systems](#), Giorgio Buttazzo, Springer, 2nd edition, 2006
 - [Real-Time Systems](#), Jane Liu, Prentice Hall, 2000
 - [An Introduction to Real-Time Systems](#), Raymond Buhr & Donald Bailey, Prentice Hall, 1998
 - [Programming with POSIX Threads](#), David R. Butenhof, Addison Wesley, 1997
 - [Computers as Components](#), Wayne Wolf, Morgan Kaufman, 2005

What are real-time systems?

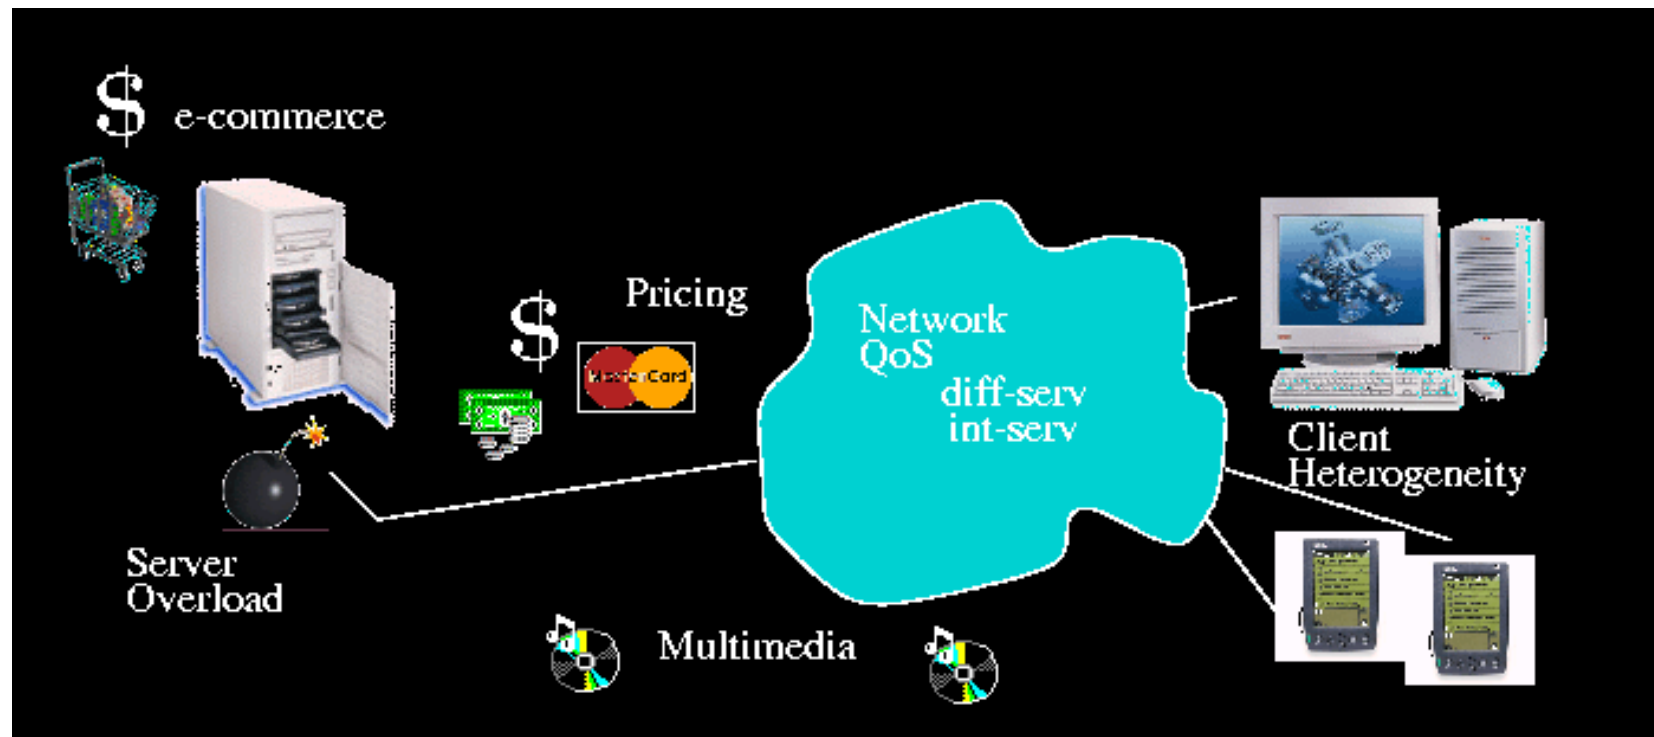
- **Study of systems where the correctness of computation depends on the timing of the results**

What are real-time systems?



Classic applications

More real-time systems



Applications of the 90s

Real-time computing today?

- Where is it going?




or

?




Where are computer systems going?

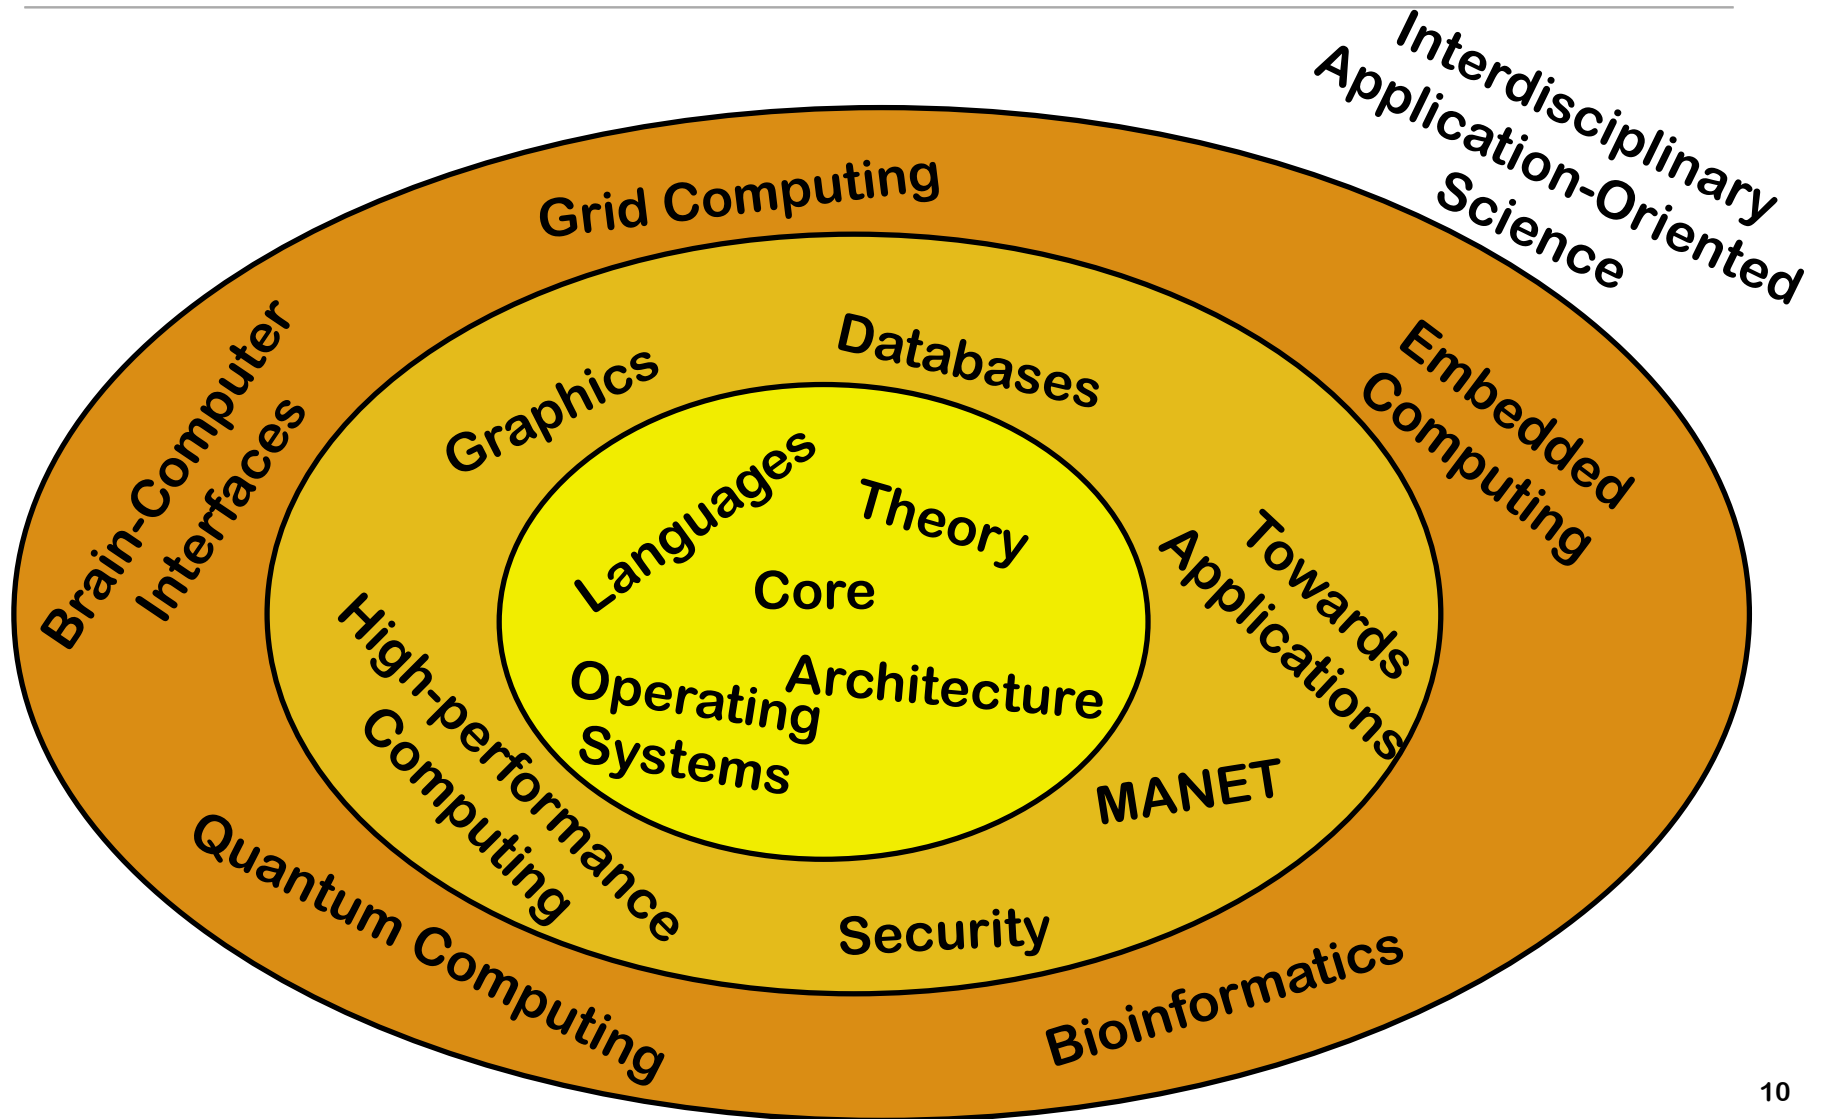
The beginning



Where are computer systems going?

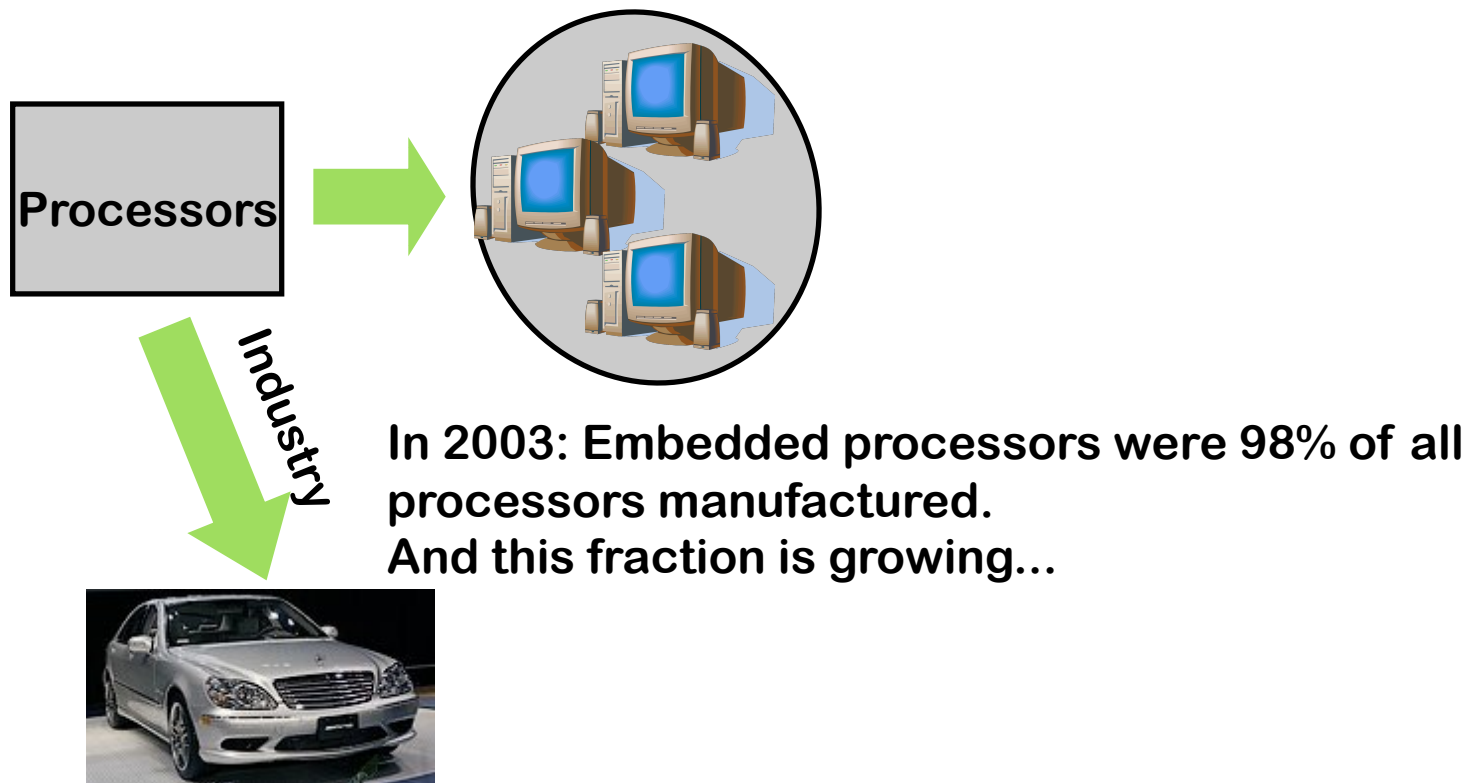


Where are computer systems going?



Real-time and embedded systems: The next frontier

Why embedded computing?



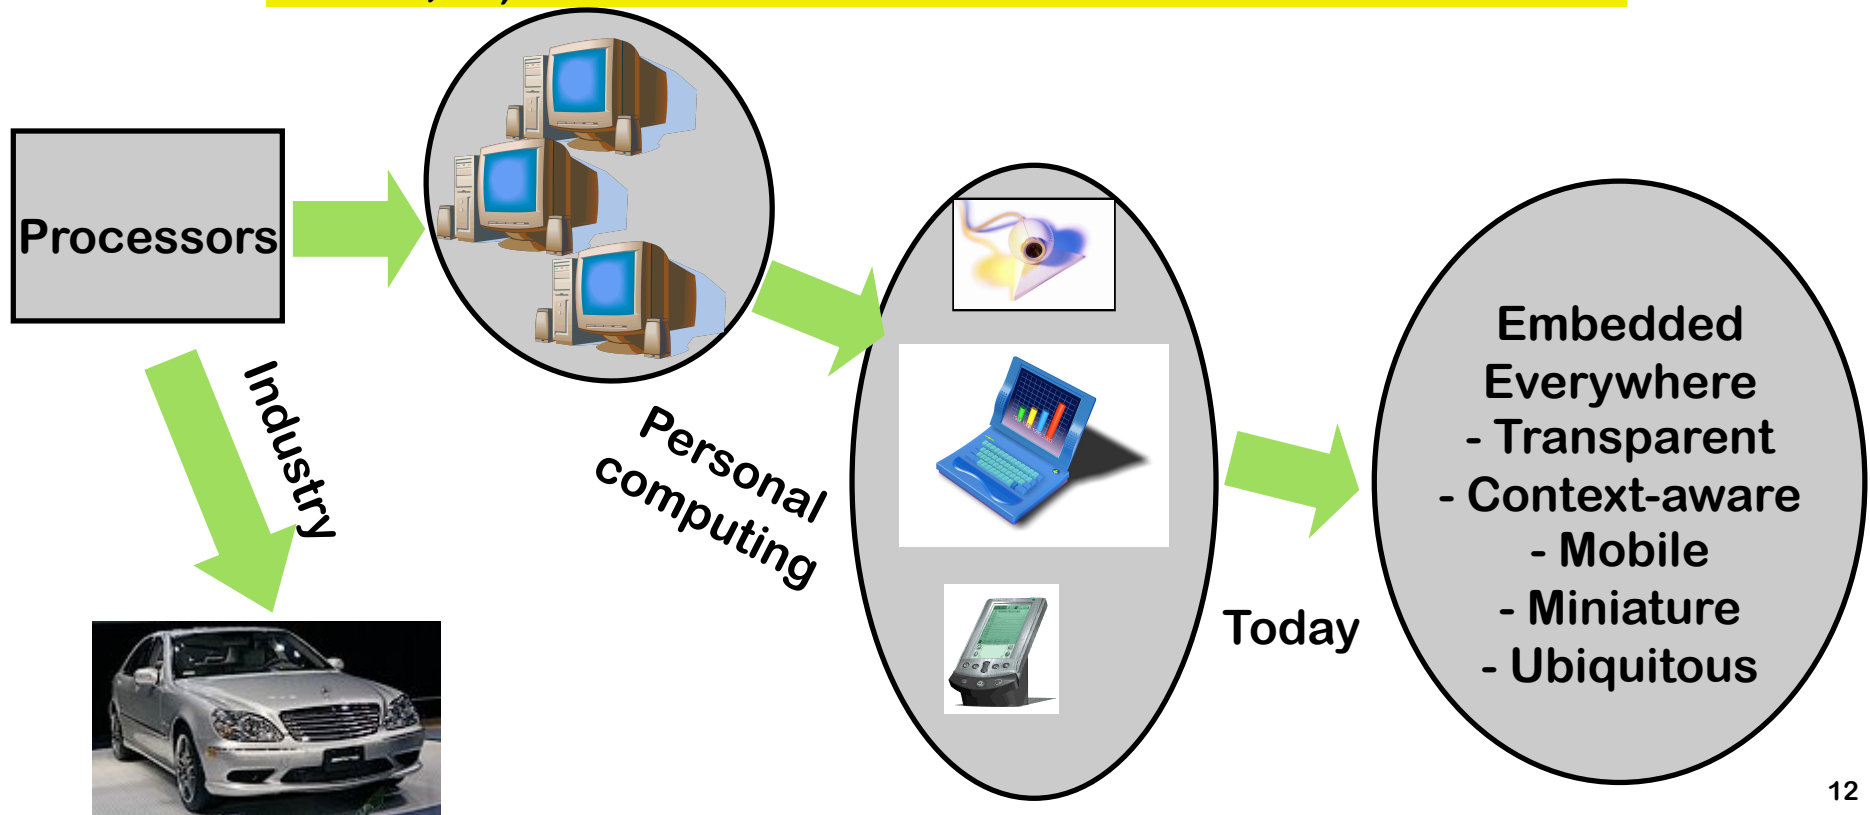
Real-time and embedded systems: The next frontier

Trend 1:

Invisible (embedded) computing, implicit interfaces

Context-aware computing (new sensors, new effectors)

Ubiquitous (instrument what we use most: attire, personal effects, ...)



Real-time and embedded systems: The next frontier

Trend 2:
Human population/input is finite;
Future data sources are increasingly
embedded devices with autonomy


Autonomic
Computing

Emergency
Response


Situation-
Aware
Distributed
Embedded
Systems

Processors

Industry



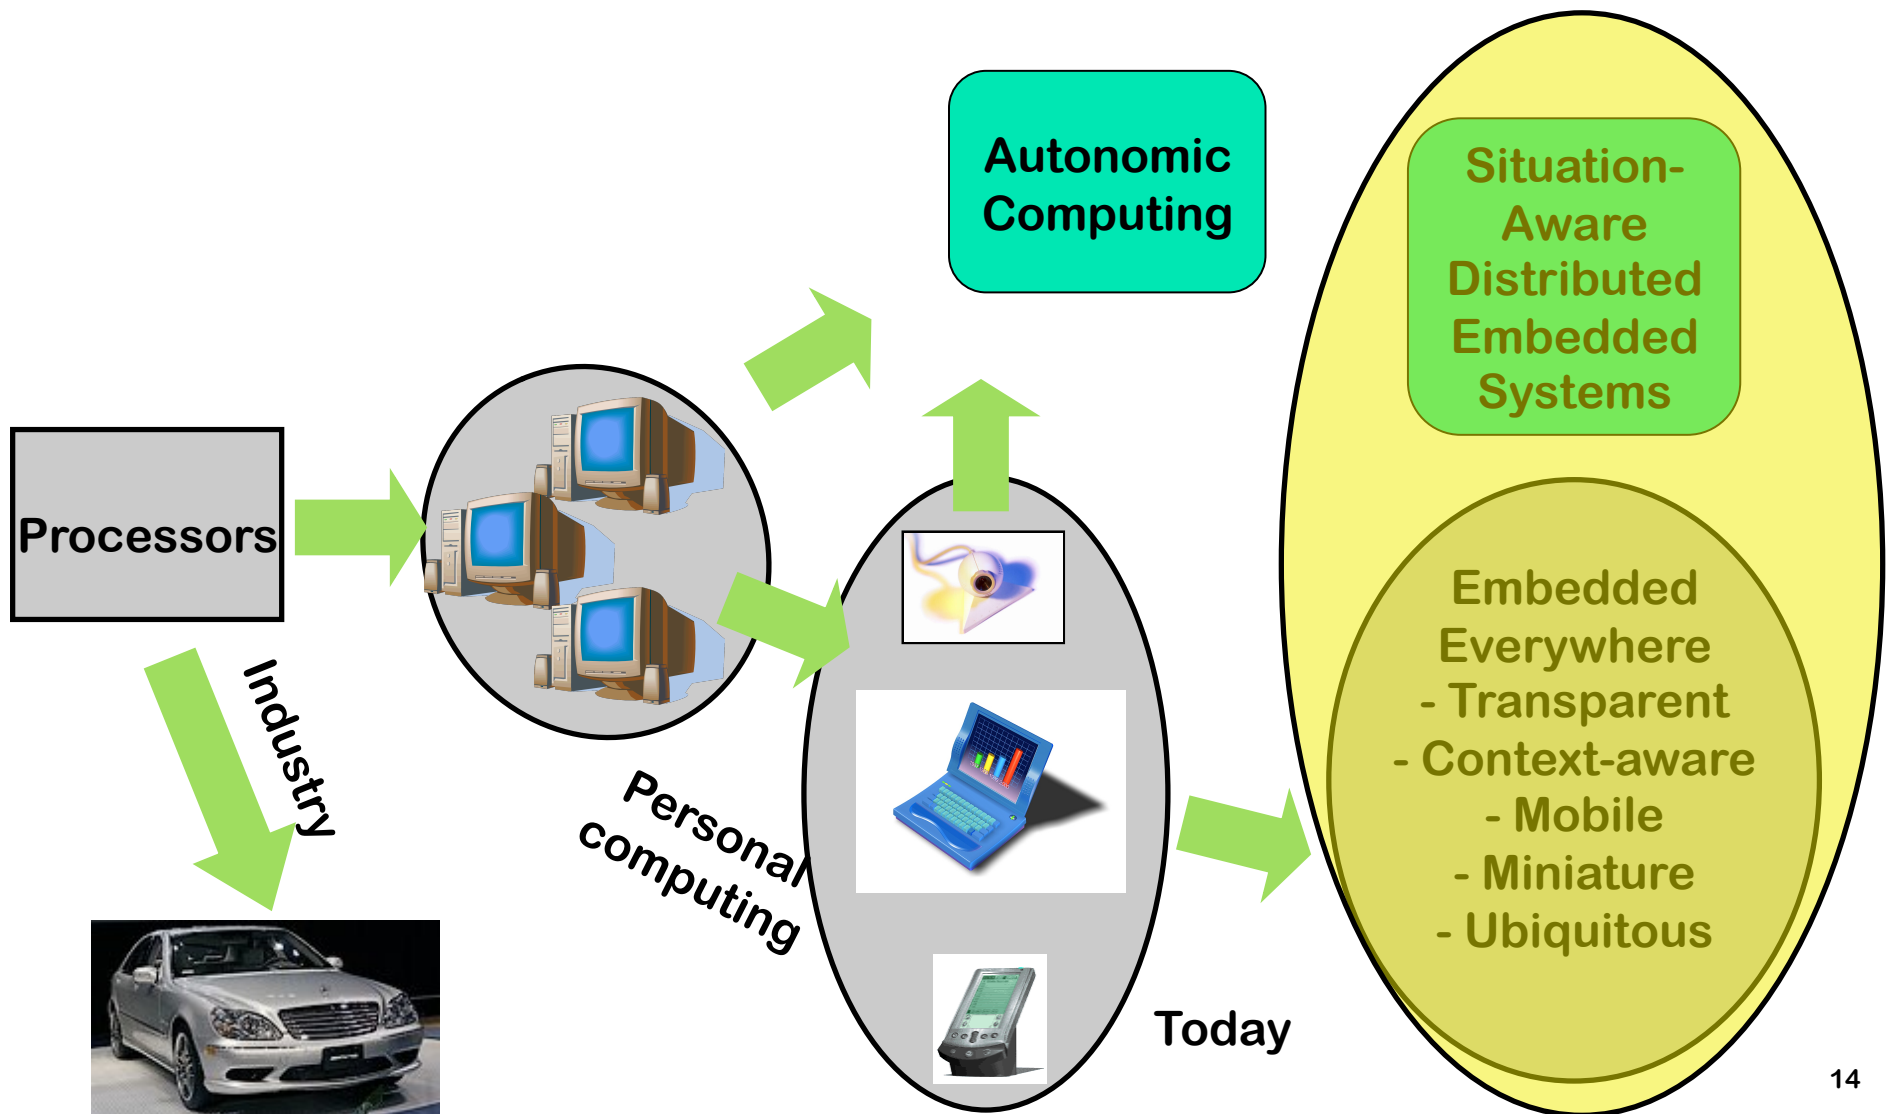
Personal
computing



Today

Embedded
Everywhere
- Transparent
- Context-aware
- Mobile
- Miniature
- Ubiquitous

Real-time and embedded systems: The next frontier



So, what does “real-time” mean?

Fast?

Predictable?

Reliable?

All of the above?

Other?

Why predictability?



**Example: Going to the airport
Which route would you choose?**

Route 1: 15 min (\$1 Toll)

**Route 2: 5 min - 45 min, with 15 min average
(free)**

You pay for predictability.



The real-time computing roadmap

- Scheduling!
- Milestone 1 (1960s): Cyclic executive scheduling
 - Fixed task set
 - Known arrivals
 - Known resources
- Milestone 2 (1970s): Rate monotonic analysis
 - Tasks could be added or removed
- Milestone 3 (1980s): The Spring scheduling approach
- Milestone 4 (1990s): Quality of service
- Milestone 5 (2000s): Feedback-controlled scheduling



**Fewer assumptions
about workload**




The changing scope of real-time computing

- **Classical view**
- **Definition: correctness of computation depends on the time at which results are generated.**
- **Applications: embedded control systems**
 - **Perfect knowledge of resources**
 - **No conflicts of interest**
- **Assumptions: hard real-time**
 - **Meeting deadlines is very important for correctness**
 - **Deadlines are inflexible**

New real-time definitions?

- Systems with low variance in performance measurements



- Systems with guaranteed metrics in which time is either in the numerator or in the denominator
 - Response time = Service time + Queuing time
 - Utilization = Used resource units/time
 - Service throughput = Produced data units/time

New assumptions

- Unknown resource requirements
- Elastic time constraints
 - Flexible periods (if any)
 - Flexible deadlines
- Adaptation capability
 - Changing algorithm version
 - Adaptive data quality
- Steady-state versus transient metrics

What we will cover in this course

- **Basic concepts (tasks, threads, blocking, priorities, importance, resource partitioning, QoS, etc.)**
- **Periodic versus aperiodic task models and their implications**
- **Optimality results in real-time scheduling**
- **Fixed-priority and dynamic-priority aperiodic-task servers**
- **Blocking, synchronization, and resource access protocols**
- **Overload, quality of service, and system design for graceful degradation**
- **Hardware/software co-design for real-time embedded systems**
- **Engineering methodology**
- **Reliability and fault tolerance**

Course organization

- **Assignments [55%]**
 - Programming assignments: 40%
 - Problem sheets: 10%
 - One-pagers: 5% (individual work)
- **Assignments to be completed in groups**
 - Groups (of 3 or 4 students) will be assigned - by the end of this week
 - All assignments to be submitted on the due date via electronic handin
- **Examinations [40%]**
 - Final exam: 40%
- **Class participation [5%]**
 - Attendance, completion of surveys, questions in class, ...

Expected background

- Understanding of operating systems
- Interrupt handling
- Multithreaded programming
- C/C++ programming skills
- Useful habits
 - Regular reading
 - Early start on assignments
 - Learn from manuals

That's it for today!

Questions?