

Lock-free (wait-free) synchronization

EECE 494

Sathish Gopalakrishnan



- Locks can be problematic.
- How do we achieve synchronization without locks?
 - Scheduling
 - Non-blocking synchronization

Eliminating locks

- One use of locks is to coordinate multiple updates of single piece of state. How to remove locks here?
 - **Duplicate state so each instance only has a single writer.**
 - (Assumption: assignment is atomic.)
- Circular buffer
 - *Why do we need a lock for a circular buffer?*
 - To prevent loss of update to buf.n (buffer size). No other reason.
 - What is buf.n good for?
 - Signaling buf full and empty.
 - How else to check this?
 - Full: $(\text{buf.head} - \text{buf.tail}) == N$
 - Empty: $\text{buf.head} == \text{buf.tail}$
- *Can we use these facts to eliminate locks in get/put?*



Lock-free sync (1 producer/1 consumer)

```
int head = 0, tail = 0;
char buf[N];
void put(char c) {
    while((buf.head - buf.tail) == N)
        wait();
    buf.buf[buf.head % N] = c;
    buf.head++;
}

void get(void) {
    char c;
    while(buf.tail == buf.head)
        wait();
    c = buf.buf[buf.tail % N];
    buf.tail++;
    return c;
}
```

All shared variables have single writer (no lock needed):

head – producer

tail – consumer

buffer:

head != tail then no overlap and

buf[head] – producer

buf[tail] – consumer

head = tail then

empty and consumer

waits until head != tail

invariants:

not full: once not full true, can

only be changed by producer


not empty: once not empty can

only be changed by consumer



Locks vs. Explicit Scheduling

- Race condition = bad interleaving of processes.
 - We've used locks to prevent bad interleaving
 - could use scheduler to enforce legal schedules.
- Examples:
 - doctor's appointment vs. emergency room
 - dinner reservation vs. showing up
 - run processes sequentially vs. acquire locks
- Tradeoffs?



- How about getting correct interleaving by detecting and retrying when a bad interleaving occurred?
 - Don't need locks to synchronize.
- Example: `hits = hits + 1;`
 - A) Read hits into register R1.
 - B) Add 1 to R1 and store it in R2.
 - C) Atomically store R2 in hits only if `hits==R1`. (i.e. CAS)
 - If store didn't write goto (A)
- Can be extended to any data structure:
 - A) Make copy of data structure, modify copy.
 - B) Use atomic word compare-and-swap to update pointer.
 - C) Goto A if some other thread beat you to the update.

Non-blocking synchronization

- Other names:
 - Wait free synchronization, Lock free synchronization.
 - [Optimistic concurrency control](#).
 - Initial work by Leslie Lamport, Maurice Herlihy.
- Modern machine have support for it:
 - **x86 CMPXCHG, CMPXCHG8B** – Compare and Exchange.
 - Someone wrote an entire OS with no locks!
- Useful properties:
 - Synchronizes with interrupt handlers.
 - Remove overhead (CPU/memory) locks.
 - Deals with failures better (e.g., process dies with locks).
- Issues:
 - Lots of retrying under high load.



A Lock-Free Multiprocessor OS Kernel

Henry Massalin and Calton Pu
Columbia University
June 1991



Where are we?

- Just to ground ourselves: Welcome to 1991
- Previously - Synthesis V.0
 - Uniprocessor (Motorola 68020)
 - No virtual memory
- 1991 - Synthesis V.1
 - Dual 68030s, virtual memory, supports threads
 - A significant step forward
 - A fairly ground-up OS



Example: Stack code

```
Pop() {  
    retry:  
        old_SP = SP;  
        new_SP = old_SP + 1;  
        elem = *old_SP;  
        if (CAS(old_SP, new_SP, &SP) == FAIL)  
            goto retry;  
    return elem;  
}
```



Stack code (Clarification)

```
Pop() {
```

```
  retry:
```

```
    old_SP = SP;
```

```
    new_SP = old_SP + 1;
```

```
    elem = *old_SP;
```

```
    if (CAS(old_SP, new_SP, &SP) == FAIL)
```

```
        goto retry;
```

```
    return elem;
```

```
}
```

local variables – can't change underneath us

“global” (at least to the data structure) ...other threads can mutate this at any time



Stack code (Stages)

```
Pop() {  
  retry:
```

- ★ 1. Write down preconditions
- ★ 2. Do computation
- ★ 3. Commit results (or retry)

```
★ old_SP = SP;
```

```
↕ new_SP = old_SP + 1;
```

```
↕ elem = *old_SP;
```

```
↕ if (CAS(old_SP, new_SP, &SP) == FAIL)
```

```
↕   goto retry;
```

```
  return elem;
```

```
}
```



- Specifically in the work by Massalin and Pu:
 - Saved state is only one or two words
 - Commit is done via Compare-and-Swap (CAS) or Double Compare-and-Swap (DCAS)
- Only two words?
 - They claimed that all OS synchronization problems can be solved while needing to touch only two words at a time
 - Impressive?

- Build data structures that work concurrently.
 - Stacks, queues, linked lists, ...
- Then, build the OS around these data structures.
- If all else fails:
 - Create a single “server” thread for the task
 - Callers then...
 1. Pack the requested operation into a message
 2. Send it to the server (using lock-free queues)
 3. Wait for response/callback/...
 - The queue effectively serializes operations.

Impact

- This paper spawned a lot of research on DCAS.
 - Improved lock-free algorithms
 - The utility of DCAS
- DCAS is not supported on modern hardware.
 - “DCAS is not a silver bullet for non-blocking algorithm design.”



What should you know?

- What is lock-free (wait-free) synchronization?
- How can it be achieved?
- What special instructions are needed?
- Read for better background:
 - **Wait-Free Synchronization**, Maurice Herlihy. ACM Transactions on Programming Languages and Systems, Vol. 13, No. 1, January 1991 (pages 124–149).
 - **A Lock-Free Multiprocessor OS Kernel**, Henry Massalin and Calton Pu. Operating Systems Review, Vol. 26, No. 2, 1992.
 - **Real-Time Computing with Lock-Free Shared Objects**, James H. Anderson, Srikanth Ramamurthy and Kevin Jeffay. IEEE Real-Time Systems Symposium, December 1995 (pages 28–37).



Conclusions

- Optimistic synchronization is effective...
 - when contention is low
 - when critical sections are small
 - as locking costs go up
- It's possible to build an entire OS without locks.
- Optimistic techniques are still applicable
 - though the implementation (DCAS) is not.

